

Lecture 16: Asset Pricing in High Dimensions

Raul Riva

FGV EPGE

March, 2026

Intro

- Last lecture: shrinkage and double selection
- Today: apply those tools to the factor zoo
- Kozak et al. ([2020](#)) ask how to estimate an SDF in a huge characteristic space
- Feng et al. ([2020](#)) ask whether one new factor still matters after the whole zoo
- Main theme: high dimension changes both estimation and inference

Why This Lecture?

- Lecture 14: weak test assets and multiple testing made the zoo look dangerous ([Lewellen et al. 2010](#); [Harvey et al. 2016](#))
- Lecture 15: regularization and valid post-selection inference gave us new tools ([Belloni et al. 2014](#))
- Today: two modern answers to the same broad problem
- One estimates a robust high-dimensional SDF
- One tests a candidate factor against a high-dimensional benchmark

- Part I: why Kozak et al. ([2020](#)) reject characteristics-sparse factor models
- Part II: how Feng et al. ([2020](#)) test whether a new factor adds marginal pricing power
- End: compare the two objects, the two methods, and the two contributions

Part I: Shrinking the Cross-Section

The Question

- Can a few characteristics summarize the whole cross-section?
- FF3 and FF5 made that look plausible ([Fama and French 1993, 2015](#))
- Lecture 14 already made us suspicious ([Lewellen et al. 2010](#); [Harvey et al. 2016](#))
- Kozak et al. ([2020](#)) ask: maybe the SDF is not sparse in characteristics at all

Why Sparse Characteristics Are a Strong Claim

- The present-value identity connects observed prices with expected profitability and investment
- Many characteristics can proxy for profitability, investment, distress, ...
- We don't actually observe expectations
- Sparsity in characteristics means extreme redundancy across anomalies

- Start with a $N \times 1$ vector of excess returns R_t
- Z_t : $N \times H$ matrix of firm characteristics
- $F_t = Z_{t-1}^\top R_t$ are factor portfolios

They consider an SDF of the following form:

$$M_t = 1 - b_{t-1}^\top (R_t - \mathbb{E}_{t-1}[R_t])$$

where $b_{t-1} = Z_{t-1}^\top b$

- If a given $b_j = 0$, then the j th characteristic does not enter the SDF at all
- Sparsity in the SDF is equivalent to sparsity of b

Plugging one equation into the other, we can rewrite the SDF as

$$M_t = 1 - b^\top (F_t - \mathbb{E}_{t-1}[F_t]),$$

Plugging one equation into the other, we can rewrite the SDF as

$$M_t = 1 - b^\top (F_t - \mathbb{E}_{t-1}[F_t]),$$

- A conditional model for the SDF entails $\mathbb{E}_{t-1}[M_t R_t] = 0$
- Then it should be the case that

$$\mathbb{E}[M_t F_t] = 0$$

since factors are themselves portfolios of returns!

Plugging one equation into the other, we can rewrite the SDF as

$$M_t = 1 - b^\top (F_t - \mathbb{E}_{t-1}[F_t]),$$

- A conditional model for the SDF entails $\mathbb{E}_{t-1}[M_t R_t] = 0$

- Then it should be the case that

$$\mathbb{E}[M_t F_t] = 0$$

since factors are themselves portfolios of returns!

- Solving for b gives:

$$b = \Sigma^{-1} \mu$$

where $\Sigma = \mathbb{E}[(F_t - \mathbb{E}[F_t])(F_t - \mathbb{E}[F_t])^\top]$ and $\mu = \mathbb{E}[F_t]$

Naive Estimation Fails

$$\hat{b}^{naive} = \hat{\Sigma}_1^{-1} \bar{\mu}$$

- This is just the sample analog of the exact formula
- Problem 1: expected returns are noisy
- Problem 2: H can be large relative to T , screwing up the covariance estimate
- Result: the estimator overfits means and loads too much on weak directions

Rotate to Principal Components

$$\Sigma_1 = QDQ^\top, \quad P_t = Q^\top F_t, \quad b_P = D^{-1}\mathbb{E}[P_t]$$

- PCs rank directions by variance
- Low-eigenvalue PCs are statistically fragile
- Directions of low variance generate huge coefficients unless the premia are estimated to be zero
- If we mis-handle them, the estimated SDF becomes wild

The Economic Prior

- Suppose you know Σ
- They assume a Gaussian prior for μ :

$$\mu \sim \mathcal{N}\left(0, \kappa^2 \frac{\Sigma^\eta}{\tau}\right), \quad \tau = \text{tr}[\Sigma] = \sum_{j=1}^H d_j$$

where d_j is the j -th eigenvalue of Σ

- Recall that $\Sigma^\eta = Q D^\eta Q^\top$
- If η is large, directions of low eigenvalues are taken to zero a priori

The Economic Prior

- Suppose you know Σ
- They assume a Gaussian prior for μ :

$$\mu \sim \mathcal{N}\left(0, \kappa^2 \frac{\Sigma^\eta}{\tau}\right), \quad \tau = \text{tr}[\Sigma] = \sum_{j=1}^H d_j$$

where d_j is the j -th eigenvalue of Σ

- Recall that $\Sigma^\eta = Q D^\eta Q^\top$
- If η is large, directions of low eigenvalues are taken to zero a priori
- Notice that $b \sim \mathcal{N}(0, \kappa^2 \frac{\Sigma^{\eta-2}}{\tau})$
- The variance about b is huge if η is too small
- They pick $\eta = 2$

The Shrinkage Estimator

- We get a spherical prior on coefficients: $b \sim \mathcal{N}\left(0, \frac{\kappa^2}{\tau} I_H\right)$
- Now they assume factors come from an i.i.d. Gaussian distribution
- This lets you derive the posterior distribution of b given the *observed* mean of factors:

$$\hat{b} = (\Sigma + \gamma I_H)^{-1} \bar{\mu}$$

where $\gamma = \frac{\tau}{T \cdot \kappa^2}$

- This looks like something you saw last class, right?

$$\hat{b} = \arg \min_b \{ (\bar{\mu} - \Sigma b)^\top \Sigma^{-1} (\bar{\mu} - \Sigma b) + \gamma b^\top b \}$$

- That estimator is equivalent to the argmin of the above objective function
- Recall that $\|b\|_2^2 = b^\top b$
- This is like Ridge but with a non-standard loss function

Penalized Interpretation

$$\hat{b} = \arg \min_b \{ (\bar{\mu} - \Sigma b)^\top \Sigma^{-1} (\bar{\mu} - \Sigma b) + \gamma b^\top b \}$$

- That estimator is equivalent to the argmin of the above objective function
- Recall that $\|b\|_2^2 = b^\top b$
- This is like Ridge but with a non-standard loss function
- Another equivalent optimization problem is

$$\hat{b} = \arg \min_b \{ (\bar{\mu} - \Sigma b)^\top (\bar{\mu} - \Sigma b) + \gamma b^\top \Sigma b \}$$

- This is like Ridge but with a different penalty
- $b^\top \Sigma b$ is the maximal Sharpe ratio attained. Why?

Why Add an ℓ_1 Penalty?

$$\hat{b} = \arg \min_b \left\{ (\bar{\mu} - \Sigma_1 b)^\top \Sigma^{-1} (\bar{\mu} - \Sigma_1 b) + \gamma_2 b^\top b + \gamma_1 \sum_{j=1}^H |b_j| \right\}$$

- ℓ_1 can drop useless factors
- But pure lasso is bad with correlated predictors
- And characteristic portfolios are very correlated
- This is like the Elastic Net, with a non-standard loss function
- Their main empirical message: ℓ_2 matters much more than sparse selection


How To Pick Penalty Parameters?

- They use K -fold cross-validation
- They use $K = 3$ and split the data in a contiguous fashion
- After estimating the SDF on $K - 1$ folds, they compute the out-of-sample R^2 on the holdout fold


$$R_{OOS}^2 = 1 - \frac{(\bar{\mu}_K - \bar{\Sigma}_K \hat{b})^\top (\bar{\mu}_K - \bar{\Sigma}_K \hat{b})}{\bar{\mu}_K^\top \bar{\mu}_K}$$

- Keep rotating and average the values

FF25 as a Sanity Check




(a) Raw Fama-French 25 portfolios




(b) PCs of Fama-French 25 portfolios

50 Anomalies From The Literature




(a) Raw 50 anomaly portfolios




(b) PCs of 50 anomaly portfolios

WRDS Financial Ratios: 70 Portfolios




(a) Raw WFR portfolios




(b) PCs of WFR portfolios

Considering Non-Linearities: Characteristics and Their Powers



(a) 50 anomalies



(b) WFR portfolios

- Not another factor paper: it is a paper about how to estimate the SDF in high dimensions
- It rejects the hope that a tiny set of characteristics is enough
- It gives a formal role to shrinkage in empirical asset pricing
- One line summary: the SDF is not sparse in characteristics, but it is still shrinkable
- Their view: Asset Pricing is about many *weak* factors...

Questions?

Part II: Taming the Factor Zoo

The Question

- Suppose someone proposes a new factor today
- How should we test it against the whole existing zoo?
- A new factor should bring new pricing information, not just be correlated with the old factors
- FF3 or FF5 are convenient benchmarks, but also arbitrary
- Feng et al. ([2020](#)) propose a conservative test on purpose

$$m_t = \gamma_0^{-1} (1 - \lambda_g^\top g_t - \lambda_h^\top h_t)$$

$$\mathbb{E}[r_t] = \iota_n \gamma_0 + C_g \lambda_g + C_h \lambda_h$$

- γ_0 : zero-beta rate (ι_n is a vector of ones)
- g_t : $d \times 1$ factors we want to test
- h_t : $p \times 1$ potentially huge set of control factors
- C_g and C_h : $n \times d$ and $n \times p$ matrices of covariances
- The goal is inference on λ_g
- The problem: h_t is large and we don't know which controls to include
- If p was small, how would you estimate λ_g ?

Why Standard Two-Pass Breaks

- If p is small, ordinary two-pass methods are fine
- But in the zoo, p can be comparable to n (number of assets) or T (number of periods)
- Using all factors is infeasible or very noisy
- Using FF3 or FF5 alone creates omitted-variable bias

Sparsity as a Working Assumption

- They assume the relevant control set is approximately low-dimensional
- This is the same as saying that λ_h has many zeros
- This is the same broad “bet on sparsity” logic as Lecture 15
- This is the traditional Fama-MacBeth logic but with “endogenous” controls
- The benchmark must be selected from a large set of candidates

Sparsity as a Working Assumption

- They assume the relevant control set is approximately low-dimensional
- This is the same as saying that λ_h has many zeros
- This is the same broad “bet on sparsity” logic as Lecture 15
- This is the traditional Fama-MacBeth logic but with “endogenous” controls
- The benchmark must be selected from a large set of candidates
- Key contribution: adapting the setup from Belloni et al. ([2014](#)) to the factor zoo
- Highly non-trivial: panel data *and* we do not directly observe covariances

Double Selection: Step 1

$$\min_{\gamma, \lambda} \left\{ n^{-1} \left\| \bar{r} - \iota_n \gamma - \hat{C}_h \lambda \right\|_2^2 + \tau_0 n^{-1} \|\lambda\|_1 \right\}$$

- First LASSO selects controls that help explain the cross-section of average returns
- This is the “pricing ability” margin
- It looks like the natural factor-selection step
- But it is not enough for valid inference on λ_g

Double Selection: Step 2

$$\min_{\xi_j, \chi_j} \left\{ n^{-1} \left\| \hat{C}_{g, \cdot j} - \iota_n \xi_j - \hat{C}_h \chi_j^\top \right\|_2^2 + \tau_j n^{-1} \|\chi_j\|_1 \right\}$$

- Second LASSO searches for controls whose exposures line up with exposures to g_t
- This is the omitted-variable-bias margin
- It keeps factors that may look weak for pricing but are dangerous to omit
- Notice that we have d separate LASSO problems here, one for each factor we want to test

$$(\hat{\gamma}_0, \hat{\lambda}_g, \hat{\lambda}_h) = \arg \min_{\gamma_0, \lambda_g, \lambda_h} \left\{ \left\| \bar{r} - \iota_n \gamma_0 - \hat{C}_g \lambda_g - \hat{C}_h \lambda_h \right\|_2^2 : \lambda_{h,j} = 0 \text{ for } j \notin \hat{I} \right\}$$

- Refit by OLS after taking the union of selected controls
- This removes the regularization bias from LASSO
- Same structure as post-double-selection in Lecture 15

Valid Inference on λ_g

Under some conditions (see the paper...):

$$\sqrt{T} \left(\hat{\lambda}_g - \lambda_g \right) \xrightarrow{d} \mathcal{N}_d(0, \Pi_1)$$

$$\Pi_1 = \lim_{T \rightarrow \infty} \frac{1}{T} \sum_{t=1}^T \sum_{s=1}^T \mathbb{E}[(1 - \lambda^\top v_t) (1 - \lambda^\top v_s) \Sigma_z^{-1} z_t z_s^\top \Sigma_z^{-1}], \quad \Sigma_z \equiv \text{Var}(z_t)$$

- z_t : residual from projecting g_t on the selected controls in h_t
- The asymptotic variance is long-run because dependence over time matters
- The paper also provides a Newey-West-type estimator for Π_1

Questions?


- Data ranges from 1976 to 2017
- 150 monthly factors from 1976 to 2017
- 15 standard public factors plus 135 long-short portfolios of stocks sorted on characteristics
- Test assets: 750 characteristic-sorted test portfolios
- The large test set directly answers the LNS critique ([Lewellen et al. 2010](#))

Factors Published in 2012-2016

Factor	λ_s (bp)	tstat (DS)
Cash holdings	-34	-0.42
HML Devil	54	1.04
Gross profitability	20	0.48
Organizational Capital	28	0.92
Betting Against Beta	35	1.45
Quality Minus Junk	73	2.03**
Employee growth	43	1.36
Growth in advertising	-12	-1.18
Book Asset Liquidity	40	1.07
RMW	160	4.45***
CMA	38	1.10
HXZ IA	51	2.11**
HXZ ROE	77	3.37***
Intermediary Risk Factor	112	2.21**
Convertible debt	-15	-1.36

Recursively Testing New Factors

Total Number of Discovered Factors Since 1994: 17



- It turns Lecture 15's double-selection logic into an asset-pricing test
- It replaces arbitrary benchmarks with selected but inference-valid benchmarks
- It responds directly to LNS and complements HLZ ([Lewellen et al. 2010](#); [Harvey et al. 2016](#))
- It makes factor discovery conservative on purpose

Questions?

- Lectures 12 and 13 built factor models
- Lecture 14 showed the zoo and the weakness of old testing standards
- Lecture 15 gave us shrinkage and valid post-selection inference
- Lecture 16: modern empirical asset pricing means regularization plus tough benchmarks
- A factor is not interesting because it has a story; it must survive high-dimensional comparison

Readings and References i

- Belloni, Alexandre, Victor Chernozhukov, and Christian Hansen. 2014. "Inference on Treatment Effects After Selection Among High-Dimensional Controls." *The Review of Economic Studies* 81 (2): 608–50. <https://doi.org/10.1093/restud/rdt044>.
- Fama, Eugene F., and Kenneth R. French. 1993. "Common Risk Factors in the Returns on Stocks and Bonds." *Journal of Financial Economics* 33 (1): 3–56. [https://doi.org/10.1016/0304-405X\(93\)90023-5](https://doi.org/10.1016/0304-405X(93)90023-5).
- Fama, Eugene F., and Kenneth R. French. 2015. "A Five-Factor Asset Pricing Model." *Journal of Financial Economics* 116 (1): 1–22. <https://doi.org/10.1016/j.jfineco.2014.10.010>.
- Feng, Guanhao, Stefano Giglio, and Dacheng Xiu. 2020. "Taming the Factor Zoo: A Test of New Factors." *The Journal of Finance* 75 (3): 1327–70. <https://doi.org/10.1111/jofi.12883>.
- Harvey, Campbell R., Yan Liu, and Heqing Zhu. 2016. "... And the Cross-Section of Expected Returns." *The Review of Financial Studies* 29 (1): 5–68. <https://doi.org/10.1093/rfs/hhv059>.

- Kozak, Serhiy, Stefan Nagel, and Shri Santosh. 2020. "Shrinking the Cross-Section." *Journal of Financial Economics* 135 (2): 271–92. <https://doi.org/10.1016/j.jfineco.2019.06.008>.
- Lewellen, Jonathan, Stefan Nagel, and Jay Shanken. 2010. "A Skeptical Appraisal of Asset Pricing Tests." *Journal of Financial Economics* 96 (2): 175–94. <https://doi.org/10.1016/j.jfineco.2009.09.001>.